

Why SSD developers need pynvme?

and why pynvme needs SPDK?

<https://github.com/cranechu/pynvme>

SSD Testing

SSD Testing

[back](#) Data Center SSDs



Intel® SSD D3-S4610 Series	Intel® SSD DC S3100 Series	Intel® SSD D5-P4326 Series
Intel® SSD DC S4600 Series	Intel® Optane™ SSD DC P4801X Series	Intel® SSD D5-P4320 Series
Intel® SSD D3-S4510 Series	Intel® Optane™ SSD DC P4800X Series	Intel® SSD DC P4101 Series
Intel® SSD DC S4500 Series	Intel® SSD DC P4618 Series	Intel® SSD DC P3700 Series
Intel® SSD DC S3710 Series	Intel® SSD DC P4610 Series	Intel® SSD DC P3608 Series
Intel® SSD DC S3700 Series	Intel® SSD DC P4608 Series	Intel® SSD DC P3600 Series
Intel® SSD DC S3610 Series	Intel® SSD DC P4600 Series	Intel® SSD DC P3520 Series
Intel® SSD DC S3520 Series	Intel® SSD DC P4511 Series	Intel® SSD DC P3500 Series
Intel® SSD DC S3510 Series	Intel® SSD DC P4510 Series	Intel® SSD DC P3100 Series
Intel® SSD DC S3500 Series	Intel® SSD DC P4501 Series	Intel® Optane™ SSD DC D4800X Series
Intel® SSD DC S3320 Series	Intel® SSD DC P4500 Series	Intel® SSD DC D3700 Series
Intel® SSD DC S3110 Series	Intel® SSD D5-P4420 Series	Intel® SSD DC D3600 Series

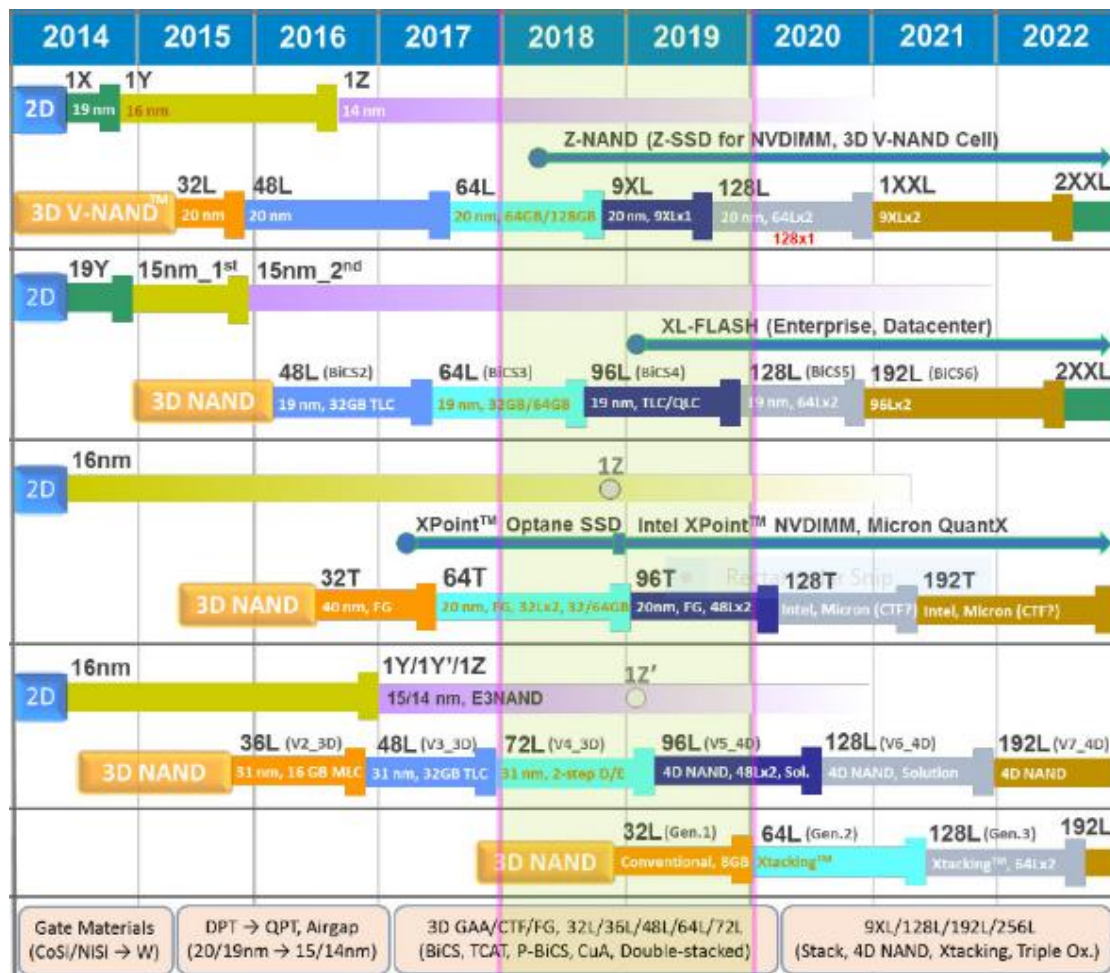
SSD Testing



SSD Testing

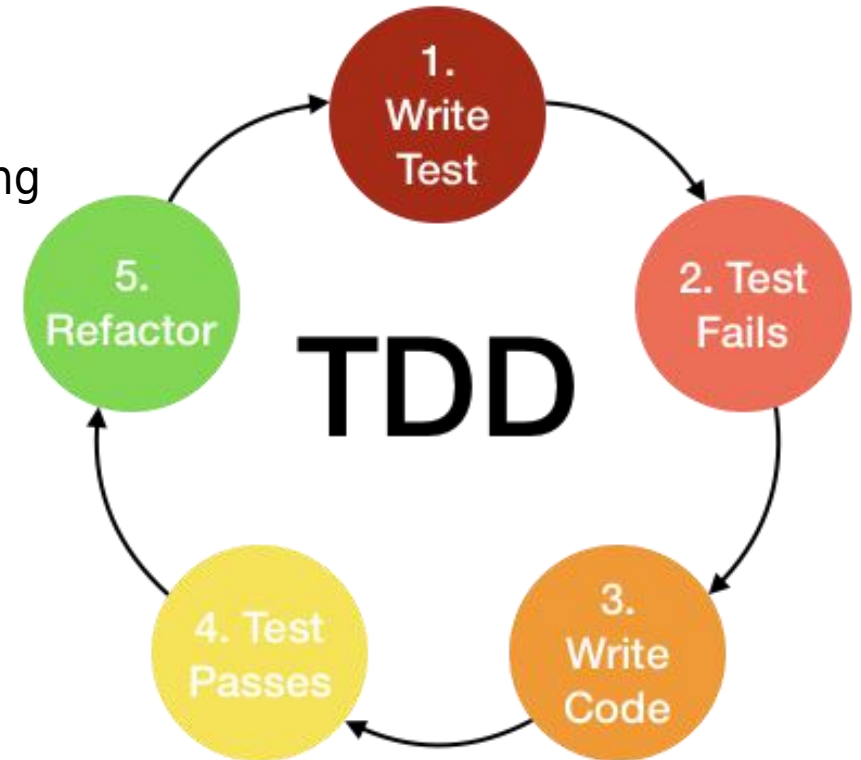
NVMe spec:

- 1.0e (Jan. 2013)
- 1.1b (July 2014)
- 1.2 (Nov. 2014)
 - 1.2a (Oct. 2015)
 - 1.2b (June 2016)
 - 1.2.1 (June 2016)
- 1.3 (May 2017)
 - 1.3a (Oct. 2017)
 - 1.3b (May 2018)
 - 1.3c (May 2018)
 - 1.3d (March 2019)
- 1.4 (June 2019)



Agile and TDD

- Challenge
 - NAND is changing
 - Applications and specifications are changing
 - Diversity on NAND and controllers
 - Continuous Qualification, testing cost
- SSD development should be ...
 - fast iterated
 - customer-oriented
 - open to change
- Waterfall v.s. Agile
 - Test-driven development (TDD)



Test Driver

- Embedded devices provide very limited resources
- We need A test-dedicated NVMe driver in host platforms:
 - exports device's features to host
 - exports device's flaws to host
 - exports device's performance to host
 - friendly to test script development
 - friendly to firmware debug
 - friendly to CI

Existed Tools

	tnvme	DM	fio	*Marks
feature	√	√	×	×
performance	×	√	√	√
scripts	√	×	√	×
debug	×	√	×	×
CI	√	×	√	×
driver	dnvme	OFA	Linux	Windows

pynvme


The `pynvme` is a python extension module.

Users can operate NVMe SSD intuitively in Python scripts. It is designed for NVMe SSD `testing` with performance considered.

Integrated with third-party tools, vscode and pytest, pynvme provides a convenient and professional solution to test NVMe devices.



Pynvme is based on SPDK/DPDK



Why SPDK?

✓ user space:

- easy for debugging
- maintainness

✓ well modularized:

- jsonrpc: for debug
- memzone: share memory between processings
- crc32

✓ best performance:

- test efficiency
- stress test

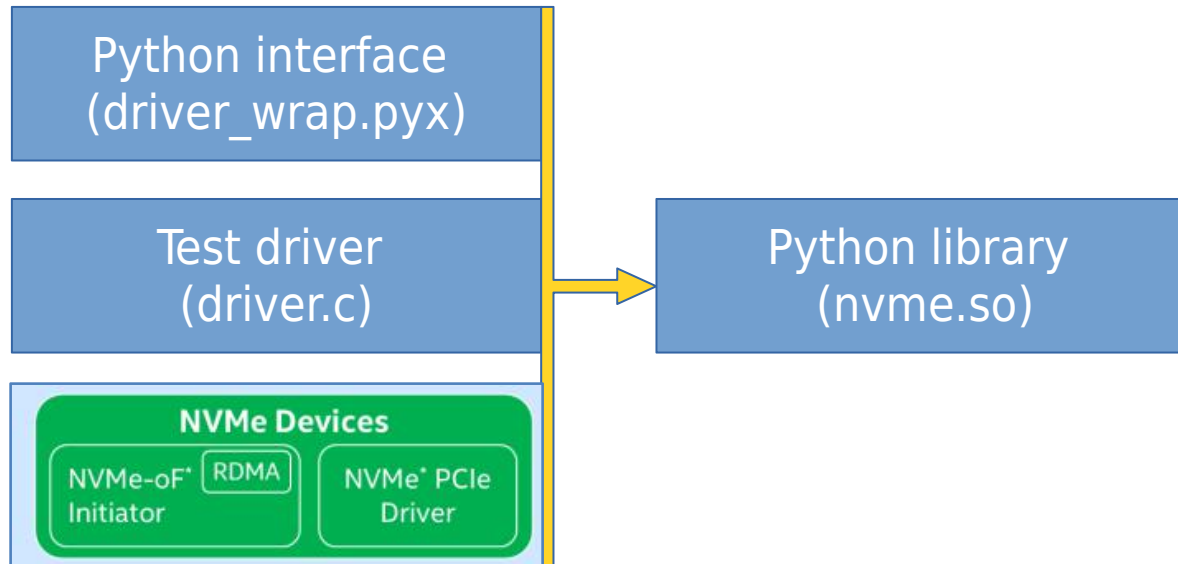
✓ open and active

- SSD, NVMe, NAND are all keeping changing!

Pynvme Architecture

Build python library with **Cython**:

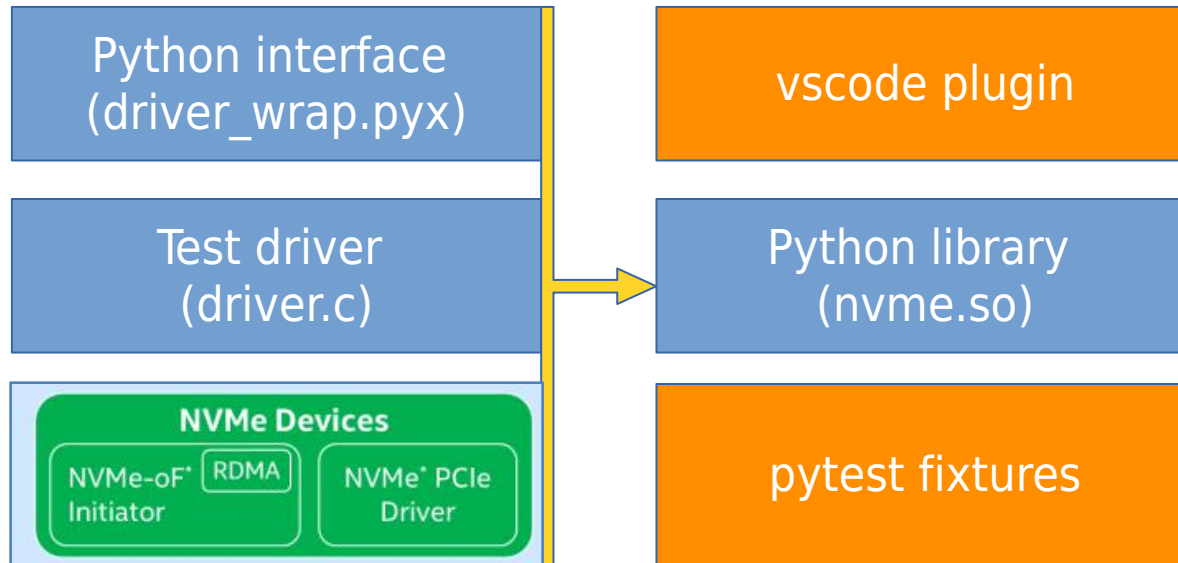
- setup.py
- driver.c
- driver.h
- cdriver.pxd
- driver_wrap.pyx
- Makefile



Pynvme Architecture

Organize test cases in **pytest**:

- mvme.so
- pytest.ini
- conftest.py
- driver_test.py



Pynvme Architecture

Python interface (driver_wrap.pyx)

Controller

Namespace

Qpair

Buffer

Test driver (driver.c)

cmdlog

checksum

MSIx

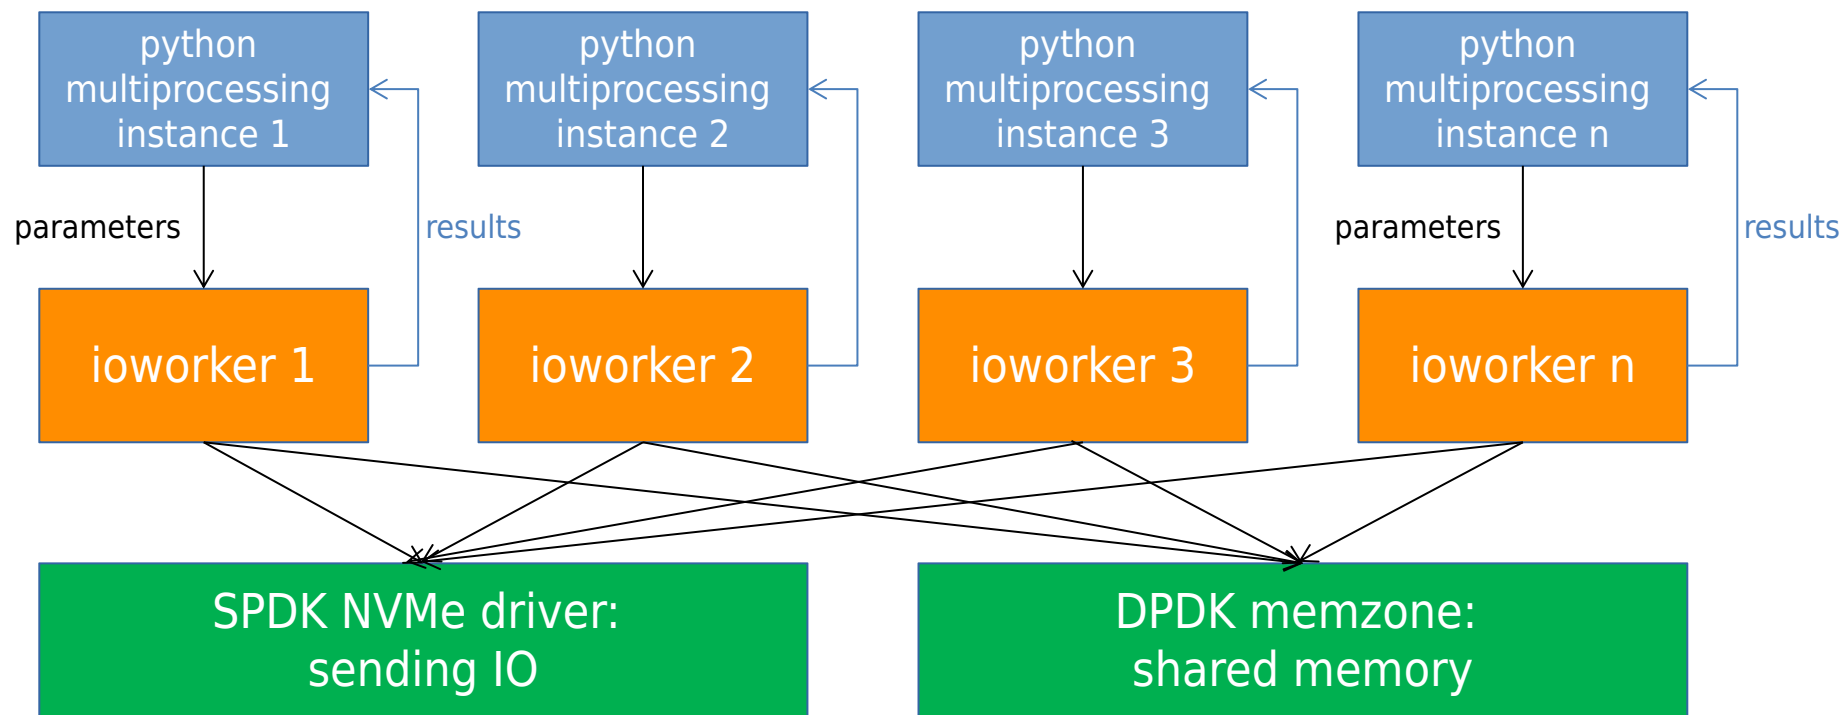
ioworker

SPDK NVMe Driver

PCIe

TCP

IOWorker



Why Python?



✓ Many beautiful and mature libraries

- Cython
- pytest
- logging
- multiprocessing
- pydoc
- os, io, time, pytemperature, statistics, yaml, json, struct, matplotlib, ...

✓ Friendly to test script development

- VSCode, Emacs, Pycharm, ...

✓ Friendly to CI: develop firmware softly

- Introducing software methodologies, processes and tools to firmware.



Pytest Execution



- `"""The pytest framework makes it easy to write small tests, yet scales to support complex functional testing for applications and libraries."""`
- `"""pytest fixtures offer dramatic improvements over the classic xUnit style of setup/teardown functions"""`
- use “make test” to start pytest sessions
 - make test
 - make test TESTS=scripts
 - make test TESTS=scripts/demo_test.py
 - make test TESTS=scripts/utility_test.py::test_download_firmware
- find test logs in test.log

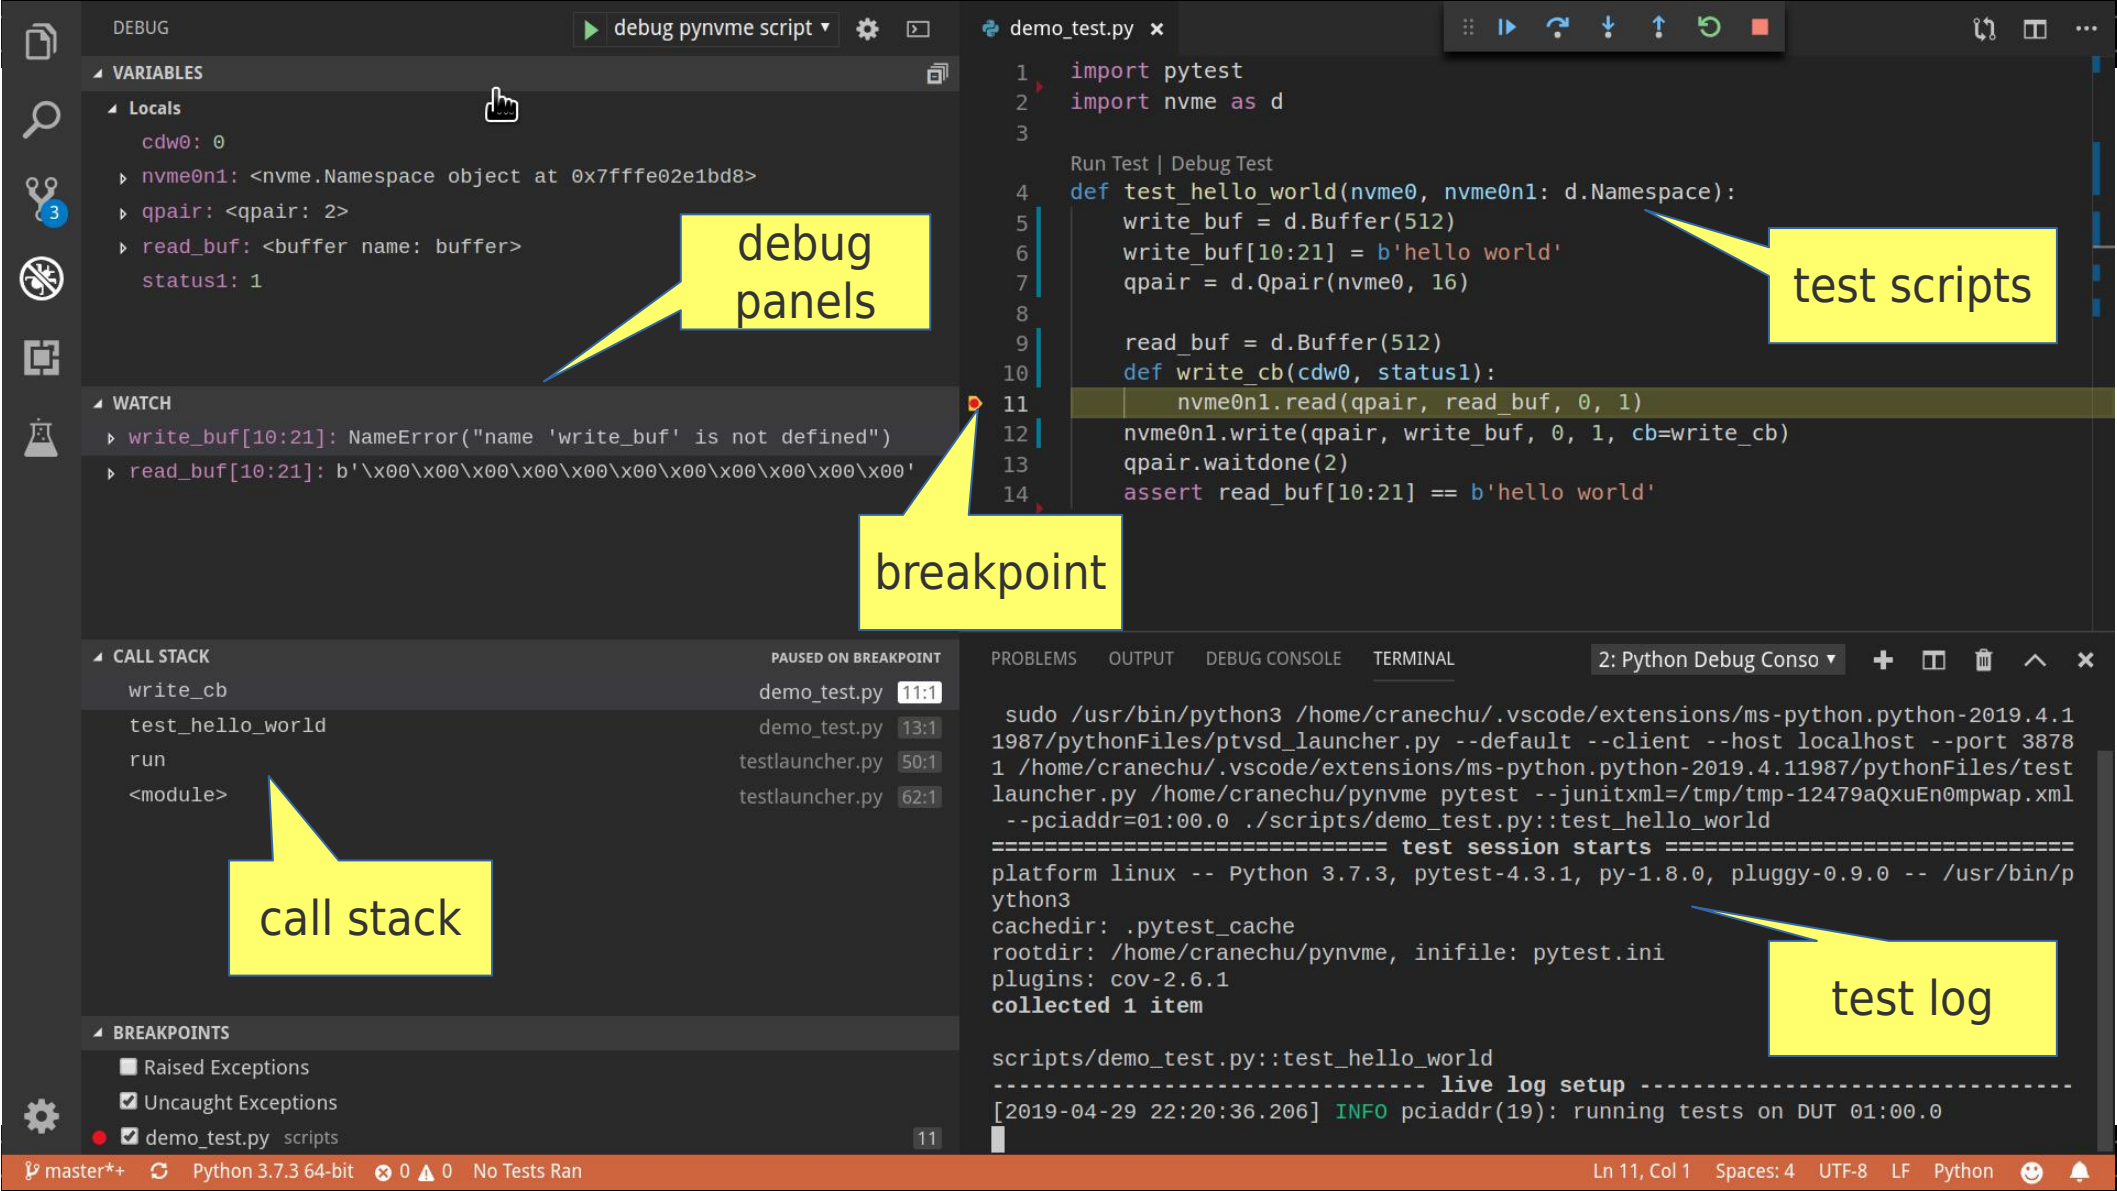
Fixtures of pynvme

- create/delete test objects. in conftest.py:
 - nvme0
 - nvme0n1
 - pcie
 - ...
- parametrize of tests
 - @pytest.mark.parametrize("qcount", [1, 2, 4, 8, 15])
 - @pytest.mark.parametrize("repeat", range(10))
- test control
 - @pytest.mark.skip("nvme over tcp")
- doc: <https://docs.pytest.org/en/latest/fixture.html>

Visual Studio Code



- VSCode is [the most popular IDE](#).
 - root user is not recommended by vscode, so users need to run sudo without a password: *sudo visudo*
- Pynvme also provides an extension to monitor device status and cmdlog in every qpair. To install the extension:
 - *code --install-extension pynvme-console-1.x.x.vsix*
- Add DUT pci address to .vscode/settings.json
 - get the BDF address with *lspci*
- *make setup; code .* # launch the vscode



debug
panels

test scripts

breakpoint

call stack

test log

test items

qpairs

```
1 import time
2 import pytest
3 import logging
4 import nvme as d
5
6
7 def test_trim_basic(nvme0, nvme0n1, verify):
8     GB = 1024*1024*1024
9     all_zeor_databuf = d.Buffer(512)
10    trimbuf = d.Buffer(4096)
11    q = d.Qpair(nvme0, 32)
12
13    # DUT info
14    logging.info("model number: %s" % nvme0.id_data(63, 24, str))
15    logging.info("firmware revision: %s" % nvme0.id_data(71, 64, str))
16
17    # write
18    logging.info("write data in 10G ~ 20G")
19    io_size = 128*1024//512
20    start_lba = 10*GB//512
21    lba_count = 10*GB//512
22    nvme0n1.ioworker(io_size = io_size,
23                    lba_align = io_size,
24                    lba_random = False,
25                    read_percentage = 0,
26                    lba_start = start_lba,
27                    io_count = lba_count//io_size,
28                    qdepth = 128).start().close()
29
30    # verify data after write, data should be modified
31    with pytest.warns(UserWarning, match="ERROR status: 02/85"):
32        nvme0n1.compare(q, all_zeor_databuf, start_lba).waitdone()
33
34    # get the empty trim time
35    trimbuf.set_dsm_range(0, 0, 0)
36    trim_cmd = nvme0n1.dsm(q, trimbuf, 1).waitdone() # first call is longer, due to
37    start_time = time.time()
38    trim_cmd = nvme0n1.dsm(q, trimbuf, 1).waitdone()
39    empty_trim_time = time.time()-start_time
40
41    # the trim time on device-side only
42    logging.info("trim the 10G data from LBA 0x%x" % start_lba)
43    trimbuf.set_dsm_range(0, start_lba, lba_count)
44    start_time = time.time()
45    trim_cmd = nvme0n1.dsm(q, trimbuf, 1).waitdone()
46    trim_time = time.time()-start_time-empty_trim_time
47    logging.info("trim bandwidth: %0.2fGB/s" % (10/trim_time))
48
49    # verify after trim
50    nvme0n1.compare(q, all_zeor_databuf, start_lba).waitdone()
```

test scripts

test log

```
===== test session starts =====
platform linux -- Python 3.7.3, pytest-4.4.0, py-1.8.0, pluggy-0.9.0 -- /usr/bin/python3
cachedir: .pytest_cache
rootdir: /home/cranechu/pynvme, inifile: pytest.ini
plugins: cov-2.6.1
collecting ... collected 1 item

scripts/test_trim_basic.py::test_trim_basic
----- live log setup -----
[2019-04-29 12:37:06.237] INFO pciaddr(19): running tests on DUT 01:00.0
----- live log call -----
[2019-04-29 12:37:10.777] INFO test_trim_basic(14): model number: SM961 NVMe SAMSUNG 1024GB
[2019-04-29 12:37:10.778] INFO test_trim_basic(15): firmware revision: CXA75D0Q
[2019-04-29 12:37:10.779] INFO test_trim_basic(18): write data in 10G ~ 20G
[2019-04-29 12:37:16.761] INFO test_trim_basic(42): trim the 10G data from LBA 0x1400000
[2019-04-29 12:37:16.763] INFO test_trim_basic(47): trim bandwidth: 16219.27GB/s
PASSED [100%]
----- live log teardown -----
[2019-04-29 12:37:16.766] INFO script(33): test duration: 5.992 sec

----- generated xml file: /tmp/tmp-5551D69216jfy7tm.xml -----
===== 1 passed in 10.65 seconds =====
===== test session starts =====
platform linux -- Python 3.7.3, pytest-4.4.0, py-1.8.0, pluggy-0.9.0 -- /usr/bin/python3
cachedir: .pytest_cache
rootdir: /home/cranechu/pynvme, inifile: pytest.ini
plugins: cov-2.6.1
collecting ... collected 1 item

scripts/test_trim_basic.py::test_trim_basic
----- live log setup -----
[2019-04-29 12:37:45.654] INFO pciaddr(19): running tests on DUT 01:00.0
----- live log call -----
```


demo_test.py x

```
1 import time
2 import pytest
3 import logging
4 import nvme as d
5 from pytemperature import k2c
6
```

Run Test | Debug Test

```
7 def test_ioworker_with_temperature(nvme0, nvme0n1):
8     smart_log = d.Buffer(512, "smart log")
9     with nvme0n1.ioworker(io_size=8, lba_align=16, lba_random=True,
10                           qdepth=16, read_percentage=0, time=30):
11         for i in range(40):
12             nvme0.getlogpage(0x02, smart_log, 512).waitdone()
13             ktemp = smart_log.data(2, 1)
14             ctemp = k2c(ktemp)
15             logging.info("temperature: %.0.2f degreeC" % ctemp)
16             time.sleep(1)
```

test scripts

define IO patterns in ioworker's parameter list, and run in a separated process

monitor temperature at the same time

test log

cmdlog

CMDLOG Q0 x

```
1 1556546728.807225: [cmd: Get Log Page]
2 0x005f0002, 0xffffffff, 0x00000000, 0x00000000
3 0x00000000, 0x00000000, 0x7439a000, 0x00000001
4 0x00000000, 0x00000000, 0x007f0002, 0x00000000
5 0x00000000, 0x00000000, 0x00000000, 0x00000000
6 1556546728.809293: [cpl: SUCCESS]
7 0x00000000, 0x00000000, 0x00000023, 0x0001005f
8
9 1556546727.804234: [cmd: Get Log Page]
10 0x005f0002, 0xffffffff, 0x00000000, 0x00000000
11 0x00000000, 0x00000000, 0x7439a000, 0x00000001
12 0x00000000, 0x00000000, 0x007f0002, 0x00000000
13 0x00000000, 0x00000000, 0x00000000, 0x00000000
14 1556546727.806340: [cpl: SUCCESS]
15 0x00000000, 0x00000000, 0x00000022, 0x0001005f
16
17 1556546726.801221: [cmd: Get Log Page]
18 0x005f0002, 0xffffffff, 0x00000000, 0x00000000
```

PROBLEMS OUTPUT ...

Python Test Log

```
[2019-04-29 22:05:11.721] INFO test_ioworker_with_temperature(15):
temperature: 44.85 degreeC
[2019-04-29 22:05:15.756] INFO test_ioworker_with_temperature(15):
temperature: 44.85 degreeC
[2019-04-29 22:05:16.760] INFO test_ioworker_with_temperature(15):
temperature: 45.85 degreeC
[2019-04-29 22:05:18.775] INFO test_ioworker_with_temperature(15):
temperature: 45.85 degreeC
[2019-04-29 22:05:19.779] INFO test_ioworker_with_temperature(15):
temperature: 45.85 degreeC
[2019-04-29 22:05:20.782] INFO test_ioworker_with_temperature(15):
temperature: 45.85 degreeC
[2019-04-29 22:05:21.786] INFO test_ioworker_with_temperature(15):
temperature: 46.85 degreeC
[2019-04-29 22:05:22.789] INFO test_ioworker_with_temperature(15):
temperature: 46.85 degreeC
[2019-04-29 22:05:23.792] INFO test_ioworker_with_temperature(15):
temperature: 46.85 degreeC
[2019-04-29 22:05:24.795] INFO test_ioworker_with_temperature(15):
temperature: 46.85 degreeC
[2019-04-29 22:05:25.800] INFO test_ioworker_with_temperature(15):
temperature: 46.85 degreeC
[2019-04-29 22:05:26.803] INFO test_ioworker_with_temperature(15):
temperature: 46.85 degreeC
[2019-04-29 22:05:27.806] INFO test_ioworker_with_temperature(15):
temperature: 47.85 degreeC
[2019-04-29 22:05:28.809] INFO test_ioworker_with_temperature(15):
temperature: 47.85 degreeC
```

pynvme goes to 1.x

cranechu / pynvme

Watch

8

Unstar

32

Fork

15

<> Code

Issues 11

Pull requests 0

Projects 2

Wiki

Security

Insights


Settings

Branch: master

pynvme / README.md

Find file

Copy path

 cranechu change homepage video

6648db0 on Jun 3


1 contributor


1228 lines (902 sloc) | 39.1 KB


Raw

Blame

History





 pynvme

test NVMe devices in Python. [<https://github.com/cranechu/pynvme>]

pipeline


passed

license

BSD-3-Clause

release

v1.0



Thanks!



pynvme

Q & A

Live Demo